

GESER DUGAROV, Ph.D. | Software Engineer, Big Data Engineer

email: geserdugarov@gmail.com

profiles: [LinkedIn](#), [GitHub](#)

SUMMARY

Software Engineer developing core functionality on a data lakehouse platform to extract value from petabyte-scale data. Hands-on experience with Apache Spark and Apache Flink, with a focus on Hudi-Spark and Hudi-Flink integrations. [Open-source contributor](#) to Apache Hudi, focusing on streaming performance and solution usability.

Building open-source tooling for agentic software development – issue-driven orchestrators and develop-review loops that turn coding-agent CLIs into reproducible engineering pipelines. Extensive experience in research and data analysis; PhD. Strong interest in big data and distributed systems.

TECHNICAL SKILLS

Java, Python, Maven, PostgreSQL, Docker, Hadoop Ecosystem, Claude Code / Codex, Agentic workflows

WORK EXPERIENCE

May 23 – Current **Software Engineer / Big Data Engineer,**
(3+ yrs) [Huawei](#)

Development of core functionality within Huawei Cloud for big data processing on enterprise-scale clusters.

- Provided a simplified configuration system utilizing commonly used presets to overcome the complexity of managing hundreds of parameters.
- Improved performance of Flink stream writing, decreasing processing time by 2x.
- Implemented partition-level TTL, enabling customers to automate cloud storage cost management with coarse granularity.

Apr 26 – Current **Open-source: Agentic development tooling,**
(~0.5 yrs) **Personal projects**

Self-directed work on agent-augmented engineering workflows – treating coding-agent CLIs (Claude Code, Codex) as bounded executors inside reproducible pipelines.

- [agent-orchestrator](#) – GitHub-Issue-driven workflow runner. Decomposes an issue, runs an implementer agent in an isolated git worktree, opens a PR, runs an independent reviewer agent, notifies a human when ready to merge. State persists in the issue itself.
- [agent-develop-review-loop](#) – shell-only develop-review loop with HITL handoff, resumable runs, and cost telemetry parsed from agent usage metadata.

Jan 24 – Current **Apache Hudi Contributor,**
(2 yrs) **The Apache Software Foundation**

[Apache Hudi](#) is a data lakehouse platform that brings database functionality to data lakes and enables incremental processing for low-latency analytics.

- Optimized serialization and deserialization of records in Flink stream writing to Hudi table, resulting in a 30% increase in processing speed and 2x reduction in memory usage ([design doc](#), [main changes](#)). Released in [Hudi 1.0.2](#).

- Implemented 4 local optimizations ([\[1\]](#), [\[2\]](#), [\[3\]](#), [\[4\]](#)) resulting in a 10% increase in processing speed and 30% reduction in garbage collection overhead. Released in [Hudi 1.0.1](#).
- Contributed 40+ [merged pull requests](#).

Feb 22 - May 23 **Software Engineer / ML Engineer,**
(1+ yr) **Digital Research** (computer vision startup)

- Designed and implemented an event-based architecture for a system for trucks monitoring. Developed server-side image processing handling ~20,000 images per day. In production, the system reduced fleet idle time by 12%.
- Built a customer-facing web UI featuring reports and data visualizations. Also developed an internal web UI for system monitoring.

EDUCATION

PhD, Geophysics, Trofimuk Institute of Petroleum Geology and Geophysics SB RAS

MSc, Computational and Applied Mathematics, Novosibirsk State University

[Recent version of CV](#)